

## SEPTEMBER 8TH, 2015 PRE-CLASS QUESTIONS

ELLIOT CARTEE

**Question 1.** Suppose  $16N$  is significantly larger than the size of our L3 cache. What is the memory-based AI of this code? (Hint: What is the memory-based AI of just the innermost loop?)

Let us first consider the innermost loop:

```
for (k = 0; k < N; k++)  
    tmp += A[i,k] * B[k,j]
```

Here we see that we are computing a total of  $2N$  flops each time we complete this innermost loop. Since  $16N$  is much larger than our L3 cache, we are getting no reuse of the cache, and therefore are also transferring  $16N$  bytes from memory to cache every time we go through this entire loop. Since we are computing no other flops, the arithmetic intensity is:

$$(\# \text{ flops}) / (\# \text{ bytes transferred between memory and cache}) = \frac{2N}{16N} = \frac{1}{8}$$

**Question 2.** If the cache is substantially larger than  $16N$ , but substantially smaller than  $8N^2$ , then we cannot fit an entire matrix into cache, but we can fit an entire row (or column). Now in the innermost loop we are still doing  $2N$  flops, but asymptotically we only need  $N$  memory accesses to write to the cache, as  $A$  is usually stored in cache and reused, while  $B$  is not. In this case we would have an arithmetic intensity of:

$$(\# \text{ flops}) / (\# \text{ bytes transferred between memory and cache}) = \frac{2N}{8N} = \frac{1}{4}$$

**Question 3.** If the cache is big enough to hold all of  $A$ ,  $B$ , and  $C$ , then we will have to transfer  $3N^2$  double precision floating point numbers in and out of the cache, meaning that in total we are transferring  $(2 \text{ transfers}) * (8 \text{ bytes per double}) * (3N^2 \text{ doubles}) = 24N^2$  bytes transferred.

Meanwhile, we compute  $2N$  flops for each entry of  $C$ , meaning that we compute  $2N^3$  flops total, giving us an arithmetic intensity of:

$$(\# \text{ flops}) / (\# \text{ bytes transferred between memory and cache}) = \frac{2N^3}{24N^2} = \frac{N}{12}$$

**Question 4.** For a problem size of  $N$ , there are  $2N^2$  double precision floating point numbers that need to fit into the cache. Since each double is 8 bytes of memory, this means that the total memory required for problem size  $N$  is  $16N^2$  bytes.

For the L1 cache with a size of 32KB, in order to fit the entire problem into cache, we need that  $16N^2 \leq 2^{15}$ . This happens when  $N \leq \sqrt{2^{15}/16} = 45.25\dots$

For the L2 cache with a size of 256KB, in order to fit the entire problem into cache, we need that  $16N^2 \leq 2^{18}$ . This happens when  $N \leq \sqrt{2^{18}/16} = 2^7 = 128$

For the L3 cache with a size of 6MB, in order to fit the entire problem into cache, we need that  $16N^2 \leq 6 \cdot 2^{20}$ . This happens when  $N \leq \sqrt{6 \cdot 2^{20}/16} = 627.07...$  (I'm not sure of the exact number of bytes in a 6MB cache since this is not a power of 2....)

For each cache, you will have arithmetic intensity  $\frac{1}{8}$  when  $16N$  bytes is larger than the cache size, as this corresponds to the case of Problem 1 where a single row will not fit in the cache. You will get arithmetic intensity  $\frac{1}{4}$  when the cache size is between  $16N$  and  $16N^2$ , as here you can fit a row into the cache but not an entire matrix, as in Problem B. You will get arithmetic intensity  $\frac{N}{12}$  when the cache size is greater than  $24N^2$  as then you can fit both matrices into the cache, as in Problem 3.

**Question 5.** Your machine becomes CPU-bound when the arithmetic intensity is greater than the ratio between peak flop rate and peak memory bandwidth. This occurs when:

$$\text{arithmetic intensity} > \frac{\text{peak flop rate}}{\text{memory bandwidth}} = \frac{(4 \text{ cores})(16 \text{ flops per cycle})(2.4 \cdot 10^9 \text{ Cycles per second})}{25.6 \cdot 10^9 \text{ bytes per second}}$$

$$\text{arithmetic intensity} > 6 \text{ flops per byte transferred}$$

**Question 6.** Based on my above answer for Question 5, in order for this machine to become CPU-bound, the arithmetic intensity will need to be greater than 6 flops per byte. Based on my answer for Question 3, this occurs when the problem size  $N$  is bigger than 72, but all three matrices can fit in the cache, so when  $24N^2$  bytes can fit in the cache. For the 6MB L3 cache, this occurs when  $N \leq \sqrt{6 \cdot 2^{20}/24} = 512$ . So the full size range is  $72 \leq N \leq 512$ .

**Question 7.** A plot of flops per second will show an increasing line for  $N < 72$  as the machine is memory bound and the arithmetic intensity increases. Then a horizontal line for  $72 \leq N \leq 512$  where the machine is CPU-bound. Then two more flat lines corresponding to the cases of Questions 1 and 2.

A plot made using `matplotlib` follows. It was generated using the code `flops_plot.py` located in this directory

