

CS 5220

Project 1 - Matrix Multiplication

Weici Hu(wh343)
Sheroze Sherifdeen(mss385)
Qinyu Wang(qw78)

September 17, 2015

1 Introduction

In this project, we tried several methods to fine-tune square matrix multiplication. Based on the `dgemm_blocked.c`, we tried unrolling index, modifying loop sequence to take advantage of SSE, and experimenting with different optimization flags.

2 Optimization

2.1 Block Multiplication with Multiple Block Sizes

2.1.1 Approach

Working off of `dgemm_blocked.c`, we tried different block sizes to examine the performance changes.

2.1.2 Results

Figure 1 shows the performance of different approaches with various block sizes. (block sizes are a multiple of 2). The performance gain for varying block sizes is not immense but a block size of 64 performs better than other block sizes.

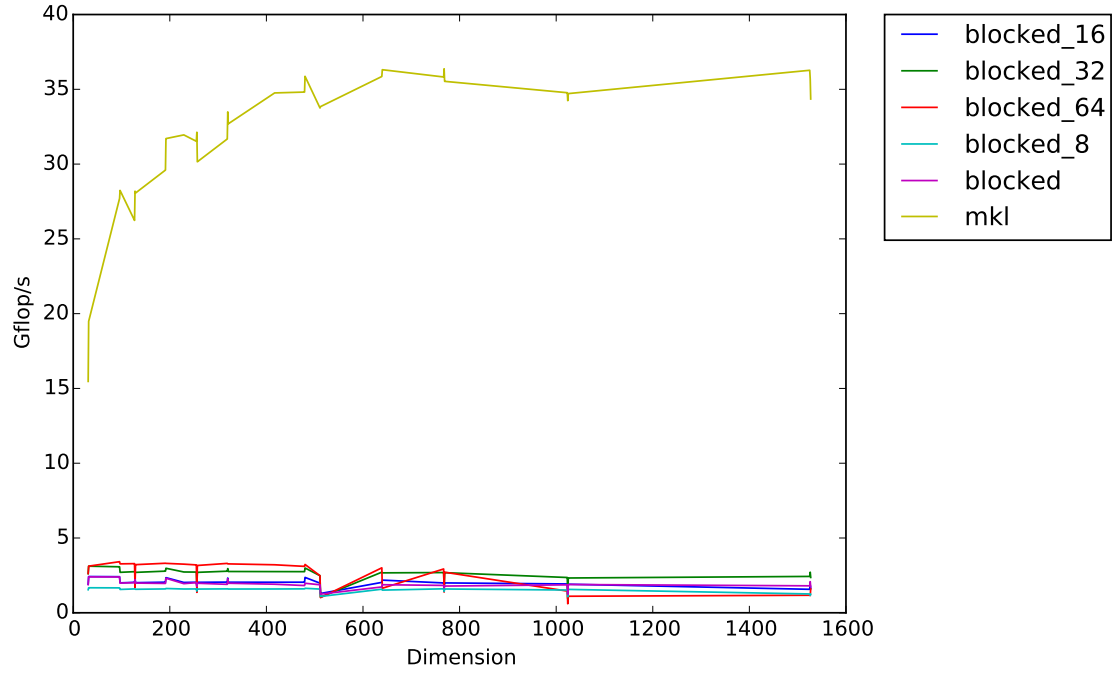


Figure 1: Block size variation

In addition, we attempted block sizes that are not a multiple of 2. Figure 2 is the comparison of the performance against a block size of 64.

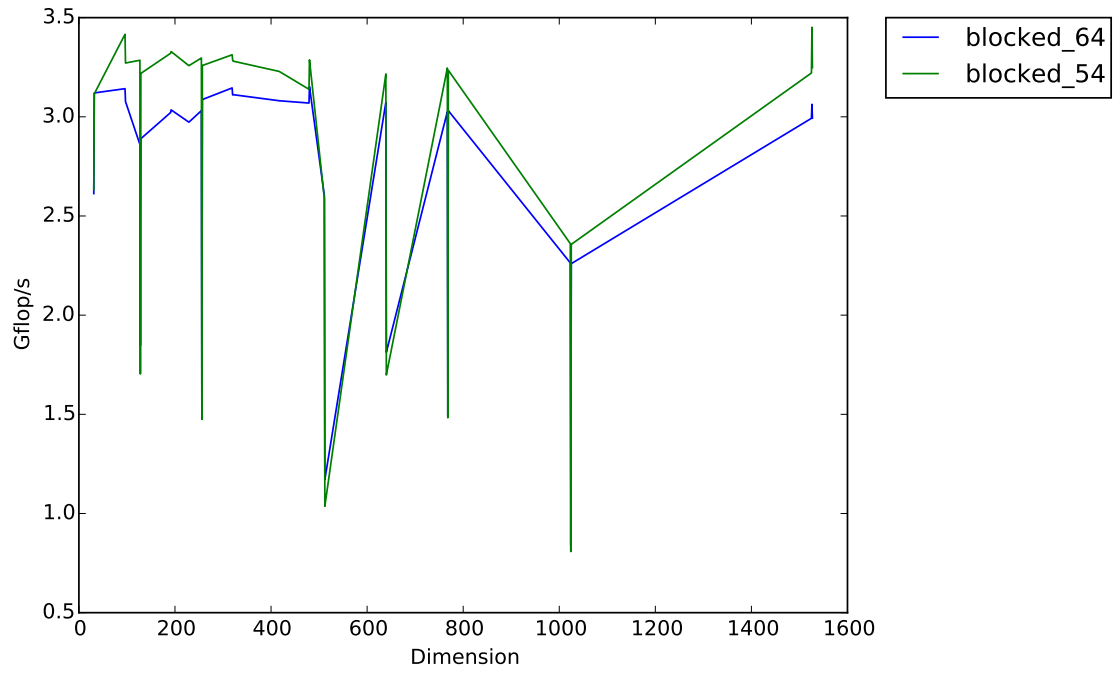


Figure 2: Block size variation

2.2 Block Multiplication with Manual Loop Unrolling

2.2.1 Approach

In this approach, we manually unrolled 4 computations in the inner most loop of the matrix multiplication of a block.

2.2.2 Results

Figure 3 compares the performance of the unrolled blocked version against the vanilla blocked approach. The unrolled versions clearly perform better than the original blocked approach but not by much.

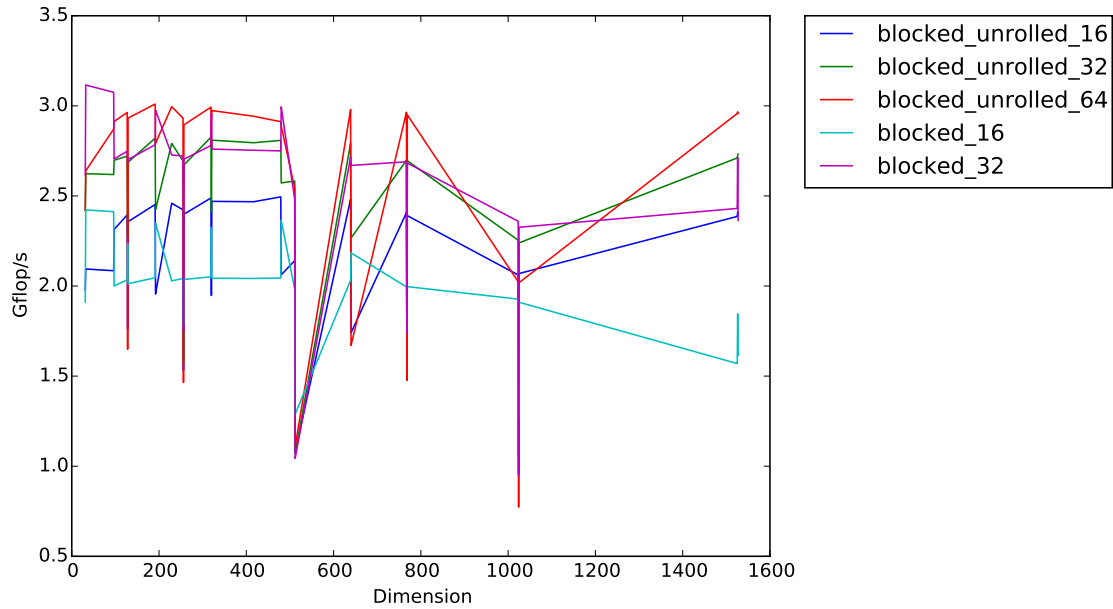


Figure 3: Unrolled blocks vs regular blocks

2.3 Loop reordering

2.3.1 Approach

2.3.2 Results

2.4 AVX Instructions

2.4.1 Approach

//Describe what we did here.

2.4.2 Results

//Add graphs here.

2.5 Compiler Optimization Flags

2.5.1 Approach

//Describe what we did here.

2.5.2 Results

//Add graphs here.

3 Next Steps

3.1 Copy Optimization